

Project Esper: Implementation of Emrakul – A DRL-Driven Immune System for Dynamic Parameter-Level Pruning in Living Neural Networks

1. Executive Summary: The Imperative for Living Architectures

The current operational paradigm of deep learning is predicated on a rigid, industrial lifecycle: training, compression, and deployment. This "train-then-prune" methodology treats neural networks as static artifacts to be optimized for a single point in time, ignoring the dynamic realities of non-stationary environments. While effective for fixed tasks, this approach is fundamentally ill-suited for "living" neural networks—systems designed for continuous, lifelong learning where data distributions shift, hardware constraints fluctuate, and new tasks emerge. **Project Esper** seeks to transcend these limitations through the implementation of **Emrakul**, a dynamic, autonomous immune system governed by Deep Reinforcement Learning (DRL).

Unlike traditional pruning, which acts as a post-hoc surgical reduction of a converged model, Emrakul operates as an intrinsic physiological process. It continuously identifies and excises "necrotic" parameters—redundant, noisy, or interfering weights—while simultaneously stimulating "synaptic neurogenesis," or the regrowth of connections, to adapt to shifting data distributions. This report provides an exhaustive theoretical and practical analysis of Emrakul's architecture, deconstructing the system into four functional quadrants:

1. **The Nervous System (State):** A high-dimensional sensory array granting the DRL agent deep visibility into the network's health via information geometry and topological embeddings.
2. **The Scalpel (Action):** A suite of differentiable actuation mechanisms, utilizing relaxation techniques like Gumbel-Softmax and synaptic flow to execute precise surgical interventions without breaking the optimization landscape.
3. **The Surgery (Protocols):** A Hierarchical Multi-Agent Reinforcement Learning (H-MARL) framework utilizing Centralized Training with Decentralized Execution (CTDE) to coordinate layer-wise optimization and resolve credit assignment.
4. **Emergent Strategies:** An analysis of the complex, adaptive behaviors—such as topological tunneling and anti-fragility—that arise from this bio-mimetic feedback loop.

Our analysis synthesizes findings from dynamic sparse training (DST), information geometry (Fisher Information), stochastic optimization, and control theory. We posit that by shifting from magnitude-based heuristics to DRL-driven policies informed by Gradient

Signal-to-Noise Ratios (GSNR) and Hessian sensitivity, Emrakul can maintain 95% sparsity while surpassing dense model performance in adaptability and generalization. This report details the theoretical underpinnings and practical implementation strategies required to realize the first truly autopoietic neural architecture.

2. Theoretical Foundations: From Static Compression to Dynamic Immunity

2.1 The Limits of Static Pruning in Continuous Learning

Traditional model compression treats neural networks as monolithic blocks of marble to be carved down to size. Techniques such as magnitude pruning, optimal brain damage, and Taylor-expansion-based sensitivity analysis effectively reduce parameter counts for inference efficiency by assuming a frozen training set and a converged model. However, within the context of Project Esper, where the "living" network must ingest continuous data streams, static pruning induces severe pathologies, most notably **catastrophic forgetting** and **plasticity loss**.

When a network is pruned based solely on past data, it loses the latent capacity—the "dormant" neurons—required to learn future, unseen concepts. Research indicates that deep reinforcement learning agents, in particular, struggle to effectively utilize network parameters, often requiring gradual magnitude pruning to maximize parameter effectiveness.¹ Yet, static schedules fail to account for the **primacy bias**, where early experiences dominate the network's topology, preventing the assimilation of new information.²

Furthermore, static pruning relies on heuristic proxies for parameter importance, most notably **magnitude pruning**, which assumes that weights with small absolute values ($|w| \approx 0$) are unimportant.⁴ While computationally efficient, this heuristic suffers from the "Magnitude Fallacy." In dynamic environments, a small weight might be a "sleeping" feature detector waiting for a specific, rare context (e.g., a "black swan" event in financial modeling or a rare edge case in autonomous driving). Pruning it based on current inactivity destroys the network's latent capability to handle future shifts.⁶

2.2 The Emrakul Doctrine: Bio-Mimetic Optimization

Emrakul addresses these limitations by treating pruning not as a reduction of size, but as a reallocation of metabolic resources—a dynamic immune response that clears "antigens" (conflicting gradients) to make room for new "antibodies" (functional circuits). This doctrine is built upon the principles of **Dynamic Sparse Training (DST)**, where the network topology evolves during training.⁷

Unlike dense-to-sparse methods that start with a heavy computational burden, DST maintains a fixed sparsity budget but dynamically rewrites connections—pruning the least effective and regrowing new ones based on gradient signals. This process mimics biological synaptic

pruning and neurogenesis. The critical innovation of Emrakul is the replacement of heuristic rules (e.g., "prune the smallest weights") with a learned DRL policy. Heuristics are rigid and local; a DRL agent, however, can optimize for long-term global rewards, balancing immediate accuracy against future plasticity. By modeling the pruning task as a Markov Decision Process (MDP)⁴, Emrakul learns to navigate the complex trade-offs between stability (retaining useful knowledge) and plasticity (freeing up capacity).

2.3 Optimization as a Bi-Level Problem

The optimization landscape of Project Esper is formally a bi-level problem.⁸

- **Lower Level:** The standard optimization of network weights θ via stochastic gradient descent (SGD) to minimize a loss function \mathcal{L} .
- **Upper Level:** The optimization of the binary mask m (topology) to maximize the agent's reward R , which is a function of the lower-level performance.

$$\max_m R(\theta^*(m), m) \quad \text{s.t.} \quad \theta^*(m) = \arg\min_{\theta} \mathcal{L}(\theta, m)$$

This interdependence creates a complex dynamic where the "Scalpel" (mask updates) directly alters the landscape of the "Nervous System" (weight updates). Emrakul's DRL agent must learn to predict how a topological change at step t will affect the trainability of the network at step $t+k$.

3. The Nervous System: High-Dimensional State Representation

For the Emrakul DRL agent to perform effective surgery, it requires a sensory system capable of perceiving the internal dynamics of the host network. A naive state representation (e.g., raw weights) is high-dimensional and opaque. The "Nervous System" of Emrakul compresses these raw signals into actionable meta-features that describe the *quality, geometry, and potential* of parameters.

3.1 Information Geometry: The Fisher Information Matrix (FIM)

A cornerstone of Emrakul's perception is the **Fisher Information Matrix (FIM)**. While magnitude tells the agent *how large* a parameter is, the FIM tells the agent *how important* that parameter is to the model's probabilistic output.

3.1.1 Diagnosing Plasticity and Primacy Bias

The FIM characterizes the local geometry of the parameter space and measures the sensitivity of the network's output distribution to changes in parameters.⁹ Formally, for a network parameterized by θ , the Fisher Information $I(\theta)$ approximates the Hessian of the loss surface.

$$\$ \$ I(\theta) = \mathbb{E}_x \sim \mathcal{D} \left\{ \right. \left. \right\}$$

High Fisher information indicates that a parameter lies in a "sharp" minimum—changing it would drastically increase error (high curvature). Low Fisher information suggests a "flat" region, where parameters are redundant or vestigial.

In Project Esper, the FIM serves as a critical diagnostic tool for **primacy bias**.² DRL systems tend to overfit to early experiences, locking parameters into configurations that are suboptimal for later tasks. The Emrakul agent monitors the trace of the FIM ($\text{Tr}(F)$). A distinct pattern in the FIM trace—specifically, a premature rigidification or saturation—signals the onset of primacy bias. By observing this state, the agent can trigger **Fisher-Guided Selective Forgetting (FGSF)**.² This protocol targets high-rigidity parameters for perturbation or pruning to restore plasticity, allowing the network to "unlearn" early overfitting without erasing generalizable knowledge. The agent learns that preserving high-Fisher weights protects "core memories," while scrubbing low-Fisher weights frees up "scratchpad" capacity.

3.1.2 Contact-Aware Fisher Information

In scenarios involving embodied agents or interaction (as implied by Project Esper's "living" nature), the Nervous System extends to **contact-aware Fisher information**.¹¹ This metric quantifies information-rich interactions, allowing the agent to prioritize parameters involved in critical decision-making boundaries (e.g., physical contact or high-stakes logical branching) over those processing background noise. By maximizing a contact-aware Fisher information measure, the agent can identify parameters that are critical for reasoning about interactions, ensuring that pruning does not compromise the agent's ability to navigate complex, interactive environments.

3.2 Gradient Dynamics: Coherence and Signal-to-Noise Ratio

While FIM measures static sensitivity, the **Gradient Signal-to-Noise Ratio (GSNR)** measures dynamic learning quality. The Emrakul agent must distinguish between weights that are actively learning and those that are oscillating due to noise.

3.2.1 Gradient Coherence as a Vital Sign

Gradient coherence refers to the alignment of gradient vectors across different training samples. When gradients align (high coherence), the network is learning a generalizable pattern. When they conflict (low coherence), the network is memorizing noise or struggling with conflicting objectives.¹²

The Emrakul agent utilizes GSNR as a state input to predict generalization performance.¹⁴

$$\$ \$ \text{GSNR}(\theta) = \frac{\mathbb{E}[\nabla_{\theta} L]^2}{\text{Var}[\nabla_{\theta} L]} \$ \$$$

A low GSNR in a specific layer indicates that the layer is failing to extract coherent features.

This serves as a signal for the agent to intervene—either by pruning the noisy connections to force the network to find a more robust path or by increasing the sparsity to regularize the layer. Research indicates that maximizing GSNR at initialization ensures better convergence¹⁴, making it a vital "prenatal" metric for determining where to regrow connections during DST cycles.

Furthermore, frozen blocks (e.g., pre-trained layers) can act as **Gradient Coherence Rectifiers**¹⁵, aligning the gradients of different samples more closely during training. Emrakul learns to leverage this by observing the GSNR profile of different layers; if a layer exhibits high coherence, it may be "frozen" or protected from pruning to act as a stable anchor, while chaotic layers are subjected to aggressive restructuring.

3.2.2 The "Flip-Flop" Detector

A specific pathology in dynamic pruning is "flip-flopping" or weight oscillation, where a parameter is pruned, regrown, and pruned again in rapid succession.¹⁶ This cycle wastes computational resources and destabilizes the network. The Nervous System tracks the sign oscillations of gradients and weight updates. High-frequency sign changes (flip-flopping) are encoded into the state vector, alerting the agent to apply a "dampening" action (e.g., temporarily freezing the weight or penalizing its modification) to enforce stability.

3.3 Topological and Graph-Based Embeddings

Neural networks are fundamentally Directed Acyclic Graphs (DAGs). Traditional state representations (vectors of statistics) lose the structural context of the network. Emrakul employs **Graph Neural Network (GNN) Embeddings** to represent the architecture itself as part of the state.¹⁸

3.3.1 Neural Graph Embedding (NGE)

Using NGE, the agent perceives the network topology not as a list of layers, but as a graph where operations are nodes and data flows are edges.¹⁹ A GNN propagates intrinsic topology information, creating a latent representation of the "cell" structure. This allows the Emrakul agent to understand dependencies—e.g., that pruning a bottleneck layer has vastly different downstream effects than pruning a redundant parallel branch. This is particularly crucial for **structured pruning** (channel/filter removal), where topological consistency must be maintained to avoid breaking the computation graph.²¹

3.3.2 Arch2Vec and Input-Output Meta-Information

To further refine this representation, Emrakul utilizes techniques like **arch2vec**, which learns joint representations of the architecture and its input-output behavior.²² By encoding the correlation between a specific sparse topology and its resulting feature map distribution, the agent can predict the functional impact of a surgical intervention before executing it. This predictive capability is essential for minimizing the "recovery time" (retraining steps) needed after a pruning event. The IO-dataset concept, using triplets of (network, input, output), allows the agent to learn a manifold of "functional equivalence," identifying disparate topologies that

yield similar performance profiles.

3.4 The Composite State Tuple

Formalizing the Nervous System, the state S_t at time t provided to the Emrakul agent is a composite tuple:

$$S_t = \langle \mathcal{E}_{top}, \mathbf{F}_{trace}, \mathbf{G}_{snr}, \mathbf{H}_{sens}, \mathbf{R}_{hist} \rangle$$

Where:

- \mathcal{E}_{top} : The GNN-based topological embedding of the current sparse architecture, capturing connectivity and bottlenecks.
- \mathbf{F}_{trace} : Trace of the Fisher Information Matrix, serving as the plasticity and memory retention metric.
- \mathbf{G}_{snr} : Layer-wise Gradient Signal-to-Noise Ratios, acting as the generalization and coherence metric.
- \mathbf{H}_{sens} : Hessian sensitivity scores, indicating robustness to quantization and perturbation.²³
- \mathbf{R}_{hist} : Historical pruning/regrowth rates, used to detect oscillatory behavior (flip-flopping).

This rich state representation allows the agent to move beyond simple magnitude pruning and engage in "informed surgery," where decisions are based on the structural, geometric, and dynamic health of the living network.

4. The Scalpel: Action Space and Surgical Mechanisms

The "Scalpel" represents the actuator mechanism of the immune system. In a standard RL setting, the action space for pruning millions of parameters is intractably large (2^N). Emrakul circumvents this "curse of dimensionality" through hierarchical decomposition, differentiable relaxation, and structured sparsity constraints.

4.1 Continuous Relaxation: The Gumbel-Softmax Mechanism

To enable end-to-end training of the pruning policy alongside the network weights, Emrakul utilizes the **Gumbel-Softmax relaxation** (also known as the Concrete distribution).²⁴ This technique bridges the gap between the discrete nature of pruning (keep/kill) and the differentiable requirements of backpropagation.

4.1.1 The Reparameterization Trick

Standard sampling from a categorical distribution (to decide whether to prune a weight) breaks the computation graph because the argmax operation has zero gradient. The Scalpel

employs the Gumbel-Max trick relaxed via a softmax function. Let $\pi_{i,j}$ be the unnormalized log-probability of keeping weight $w_{i,j}$. We sample a mask $m_{i,j}$ as:

$$m_{i,j} = \frac{\exp((\log \pi_{i,j} + g_{i,j}) / \tau)}{\sum_k \exp((\log \pi_k + g_k) / \tau)}$$

Here, $g_{i,j}$ is independent Gumbel noise sampled from $\text{Gumbel}(0, 1)$, and τ is the temperature parameter.²⁶

- **Forward Pass:** The system uses a "Straight-Through" estimator, discretizing the soft sample to a binary mask $M \in \{0, 1\}$ using argmax for the actual computation. This ensures the network sees a truly sparse architecture during inference.²⁷
- **Backward Pass:** The gradients flow through the continuous softmax approximation, allowing the DRL agent to update its policy π based on the downstream performance loss. This decoupling allows the agent to learn the *probability* of pruning a weight, rather than just the decision itself.

4.1.2 Temperature Annealing as Surgical Precision

The temperature τ controls the "softness" of the decision.²⁸

- **High τ (Exploration):** Approximates a uniform distribution. Gradients have low variance but high bias. This setting corresponds to "exploratory surgery," where the agent tests various connectivity patterns.
- **Low τ (Exploitation):** Approximates a discrete categorical distribution (argmax). Gradients have high variance but low bias. This corresponds to "precision surgery." Emrakul dynamically controls τ as a "focus" parameter. During early phases of "neurogenesis" (regrowth), high τ allows for broad exploration of connectivity. As the topology stabilizes, the agent anneals τ to 0, hardening the decisions into a fixed sparse structure.²⁴

4.2 Differentiable Subset Pruning (DSP) for Heads and Channels

For structured pruning (removing entire attention heads or filters), the Scalpel utilizes **Differentiable Subset Pruning (DSP)**.²⁹ Instead of treating each head independently, DSP views pruning as a **subset selection problem**. The goal is to select exactly K heads out of H to retain.

This is achieved using the **Gumbel-Top-K** trick.

1. The agent learns importance scores i_h for each head h .
2. These scores are perturbed with Gumbel noise: $r_h = \log(i_h) + g_h$.
3. The top K heads with the largest perturbed scores are selected.

$$h^*_1 = \text{argmax}_h \quad h^*_2 = \text{argmax}_{h \in \setminus \{h^*_1\}} r_h, \dots$$

This mechanism enforces a hard constraint on the resource budget (e.g., "keep exactly

50% of heads") while remaining fully differentiable. It allows the agent to perform combinatorial optimization within the gradient descent loop, preventing the agent from collapsing to trivial solutions (keeping all heads) or catastrophic ones (pruning everything).

4.3 "Synaptic Flow" and N:M Structured Sparsity

To ensure the sparse network can be accelerated on modern hardware (like NVIDIA Ampere/Hopper Tensor Cores), the Scalpel enforces **N:M sparsity** (e.g., 2:4 sparsity, where at most 2 non-zero values exist in every block of 4).³¹

4.3.1 Hardware-Aware Probabilistic Masking

The agent does not prune weights individually but selects *patterns*. For every block of M weights, the agent learns a prior categorical distribution and samples the N survivors without replacement.³³ If the agent selects a pattern that violates N:M sparsity, it receives a penalty. This **Linear-Space Probabilistic Framework** ensures the learned masks are directly compilable to Sparse Tensor Core instructions, achieving up to $2 \times$ theoretical speedup.³⁴

4.3.2 The Synaptic Flow Conservation

A major risk in DRL-driven pruning is the "vanishing gradient" problem for pruned weights. Once a weight is zeroed, it typically stops receiving gradients, making it impossible to revive if it becomes relevant later. Emrakul implements the Synaptic Flow protocol.⁸ The Scalpel maintains a dense set of "virtual" gradients even for pruned connections.

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta \mathcal{L}(f(x; \theta) \odot m)$$

Even if $m=0$, the underlying parameter θ is updated. If a masked weight accumulates enough magnitude (crossing a threshold defined by the DRL policy), it is unmasked. This allows for gradient-based regrowth⁷, ensuring the topology evolves toward the data distribution.

4.4 Probabilistic Masking via Hard Concrete Distributions

For more stability than standard Gumbel-Softmax, the Scalpel employs the Hard Concrete distribution.²⁵ Standard Softmax outputs are always in $(0, 1)$, never exactly 0. The Hard Concrete distribution stretches this range to (γ, ζ) with $\gamma < 0$ and $\zeta > 1$, and then clamps the values to 0 and 1 .

$$\text{mask} = \min(1, \max(0, \tilde{z}))$$

This allows probability mass to accumulate exactly at 0 and 1. This creates a "Probabilistic

"Mask" where weights have a non-zero probability of being exactly zero, reducing the variance of the gradient estimator and allowing for true sparsity during the training process itself.³⁶

5. Surgery: Protocols and Multi-Agent Coordination

Surgery in Emrakul is not a unitary event but a continuous, orchestrated process. Given the scale of deep networks, a single RL agent cannot manage every parameter. Project Esper adopts a **Hierarchical Multi-Agent Reinforcement Learning (H-MARL)** framework to coordinate the surgery.

5.1 Hierarchical Control: The "Surgeon General" and "Specialists"

Emrakul divides the control policy into two levels of abstraction³⁷:

1. **The High-Level Controller (Surgeon General):** This agent operates on a slower timescale. It observes global state metrics (validation accuracy, total sparsity, latency constraints, global GSNR) and determines the **layer-wise sparsity budgets**.
 - o *Action:* $\$A_{\text{global}} = [s_1, s_2, \dots, s_L]$, where s_i is the target sparsity for layer i .
 - o *Strategy:* It learns strategic resource allocation, e.g., that early feature extraction layers require higher density than later fully-connected layers.²¹
2. **The Low-Level Workers (Specialists):** These are layer-specific agents (or a shared policy across similar layers) that execute the specific pruning decisions to meet the budget set by the Controller.
 - o *Action:* Generating the specific Gumbel-Softmax mask M_i for layer i .
 - o *Strategy:* They optimize for local reconstruction error and gradient coherence, ensuring the layer maximally preserves information within the constraints set by the Surgeon General.

5.2 CTDE: Centralized Training, Decentralized Execution

To train this multi-agent system, Emrakul employs **Centralized Training with Decentralized Execution (CTDE)**, specifically adapting algorithms like **QMIX** or **MAPPO**.⁴⁰

- **Centralized Training (The Clinic):** During the training phase, a centralized critic network has access to the full global state S_{global} and the joint actions of all agents. It estimates the global Q-value function $Q_{\text{tot}}(s, \mathbf{a})$. In QMIX, this global Q is a non-linear combination of individual agent utilities Q_a . This allows the system to solve the **credit assignment problem**: if the model fails, the Critic can determine whether it was due to Agent A's aggressive pruning in Layer 2 or Agent B's failure to regrow in Layer 5.⁴³
- **Decentralized Execution (The Wild):** During deployment (lifelong learning), the centralized mixing network is discarded. Each layer-wise agent acts solely on its local observation o_i (local activations, local gradients). This decoupling ensures scalability and allows the immune system to operate with minimal communication overhead during

the "living" phase.⁴⁴

5.3 Reward Shaping: The Hippocratic Oath of Pruning

The reward function is the guiding moral compass of the immune system. A naive reward (just accuracy) is too sparse and delayed. Emrakul utilizes **Dense Rewards** and **Composite Objectives**.⁴

$$\text{R}_t = \underbrace{\alpha \cdot \text{Acc}_t}_{\text{Performance}} - \underbrace{\beta \cdot \text{FLOPs}_t}_{\text{Efficiency}} - \underbrace{\gamma \cdot \text{Osc}_t}_{\text{Stability}} + \underbrace{\delta \cdot \text{GSNR}_t}_{\text{Health}} - \underbrace{\epsilon \cdot \text{Spike}_t}_{\text{Trauma}}$$

- **Step-wise Feedback:** Unlike episodic RL which rewards only at the end, Emrakul provides feedback at every pruning step (e.g., change in validation accuracy), accelerating convergence by 85% compared to sparse rewards.⁴
- **Stability Penalty (Osc_t):** To prevent **flip-flopping** (the Flip-Flop Detector from the Nervous System), the agent is penalized for reversing a pruning decision made in the recent window $t-\Delta$. This incentivizes decisive, long-term actions.⁴⁶
- **Loss Spike Penalty (Spike_t):** Large sudden spikes in loss (indicating catastrophic structural damage) incur a heavy penalty. This teaches the agent to perform "minimally invasive" surgery and effectively avoid **Critical Double-Loss** scenarios where both weights and necessary data support are removed simultaneously.⁴⁷
- **Health Bonus (GSNR_t):** The agent receives a bonus for maintaining high Gradient Signal-to-Noise Ratio in the remaining connections, explicitly reinforcing the retention of coherent, generalizable features.¹⁴

5.4 Anesthesia Protocols: Managing Surgical Shock

Performing aggressive pruning on a live network can induce "shock" (instability). Emrakul implements **Anesthesia Protocols** derived from control theory and biological analogues.⁴⁹

- **Reversible Sedation:** Before a major topological restructuring (e.g., pruning a head), the agent may temporarily dampen the learning rate of the affected layers (sedation) or freeze the weights of connected layers. This prevents the propagation of "panic gradients" (high variance updates) resulting from the sudden structural change.
- **Harmonic Annealing:** A sparsity scheduler that allows periodic regrowth (waking up) to counteract excessive sparsification. By harmonically oscillating the target sparsity, the system ensures it doesn't drift into a "coma" (unrecoverable performance loss) and maintains a pool of active weights for exploration.⁴⁰

6. Emergent Strategies and Third-Order Insights

The convergence of DRL, FIM-based sensing, and Gumbel-Softmax actuation within Emrakul

gives rise to complex, emergent behaviors that extend beyond simple model compression. These strategies define the "personality" of the living network and represent the third-order insights of this analysis.

6.1 The "Lazarus" Strategy: Escaping Local Minima via Topology

Traditional optimization (SGD) struggles to escape deep local minima in the weight space. Emrakul utilizes topology as an orthogonal dimension of optimization. By temporarily degrading performance (pruning critical paths) and then aggressively regrowing them in new configurations, the agent can "tunnel" through energy barriers.⁵² This is a form of Topological Simulated Annealing.⁵³ The agent learns that a short-term loss in accuracy (high entropy action) can lead to a superior long-term basin of attraction. The system essentially "heats up" the topology (increasing Gumbel temperature τ) to detach from a suboptimal local minimum and then "cools down" (annealing τ) to settle into a more robust configuration.

6.2 Anti-Fragility and Dynamic Redundancy

Unlike static pruning which aims for the minimal effective network (efficiency), Emrakul learns to maintain dynamic redundancy (resilience). The agent preserves "ghost circuits"—sparse, inactive pathways that have high gradient potential (Synaptic Flow).⁷ These pathways act as a reserve army. When the data distribution shifts (domain drift), the agent activates these ghost circuits instantly via the Gumbel-Softmax gates. This creates an anti-fragile system that becomes more robust under stress (distributional shift) because it has maintained the capacity for adaptation, even if that capacity wasn't immediately active. This directly addresses the irreversibility problem of static pruning.

6.3 Critical Double-Loss Avoidance

A sophisticated emergent strategy is the avoidance of Critical Double-Loss.⁴⁸ In standard training, noisy data samples and redundant weights often interact destructively. Pruning weights can make the model vulnerable to specific data outliers, while data pruning (coreset selection) can make weights overfit.

The Emrakul agent, observing the FIM traces, learns to correlate specific topological configurations with data types. It develops a strategy of Synchronized Retention: it will refuse to prune weights that have high contact-aware Fisher information regarding "tail" data, even if their magnitude is low. It effectively performs "co-evolutionary" optimization of the dataset and the topology, ensuring the hardware (weights) exists to process the software (rare samples).

6.4 Temporal Feature Masking: The Breathing Network

Through Adaptive Temporal Masking (ATM)⁵⁵, the agent learns to smooth its interventions over time. Instead of binary flickering, the agent adopts a probabilistic approach where feature importance is tracked via Exponential Moving Averages (EMA).

This results in a "breathing" network where modules gently fade in and out of relevance based

on the temporal context of the data stream. During phases of high novelty, the agent expands the active parameter count (inhale). During consolidation phases, it prunes back to the most efficient core (exhale). This emergent rhythm minimizes energy consumption (FLOPs) while maximizing adaptability, mirroring biological metabolic cycles and preventing the "feature absorption" problem common in sparse autoencoders.

7. Conclusion: The Autopoietic Network

Project Esper, through the implementation of Emrakul, represents a fundamental transition from **allopoietic** AI (systems built and maintained by external creators) to **autopoietic** AI (systems that self-maintain and self-create).

The theoretical synthesis of the **Fisher Information Matrix** provides the necessary historical context to prevent amnesia; **Gumbel-Softmax** relaxation provides the surgical dexterity to manipulate topology within a differentiable framework; and **Multi-Agent RL** provides the scalable coordination required for deep architectures.

The analysis confirms that Emrakul is not merely a pruning algorithm but a metabolic engine for artificial cognition. By continuously metabolizing data to build structure, and metabolizing structure to refine information, Emrakul resolves the stability-plasticity dilemma. The future of large-scale neural networks lies not in ever-larger dense matrices, but in living, sparse substrates that evolve in real-time.

Table 1: Comparative Analysis of Pruning Paradigms

Feature	Static Pruning (Traditional)	Dynamic Sparse Training (DST)	Emrakul (Project Esper)
Trigger	Post-training / One-shot	Periodic / Schedule-based	Continuous / State-Dependent
Criterion	Magnitude (\$)	w	\$)
Scope	Global or Layer-wise (Fixed)	Layer-wise (Heuristic)	Hierarchical Multi-Agent (CTDE)
Recovery	Fine-tuning required	Gradient-based Regrowth	Synaptic Flow & Anti-Fragile Reserves
Objective	Efficiency (FLOPs/Size)	Efficiency + Accuracy	Lifelong Plasticity & Generalization
Mechanism	Hard Masking	Periodic Mask Updates	Differentiable Gumbel-Softmax Relaxation
Dynamics	Irreversible	Limited Reversibility	Topological Tunneling & Breathing

Emrakul demonstrates that the optimal topology of a neural network is not a fixed destination

but a continuous trajectory. In the hostile, non-stationary environments of the real world, the only viable strategy is not just to learn, but to evolve.

Works cited

1. In deep reinforcement learning, a pruned network is a good network - arXiv, accessed January 3, 2026, <https://arxiv.org/html/2402.12479v1>
2. Fisher-Guided Selective Forgetting: Mitigating Primacy Bias in Deep Reinforcement Learning | OpenReview, accessed January 3, 2026, <https://openreview.net/forum?id=uOGK8zwgvc>
3. Primacy Bias Through the Lens of the Fisher Information Matrix, accessed January 3, 2026, <https://fse.studenttheses.ub.rug.nl/34453/1/mAI2024MassimilianoFalzari.pdf>
4. Learning to Prune Deep Neural Networks via Reinforcement Learning - AutoML.org, accessed January 3, 2026, https://www.automl.org/wp-content/uploads/2020/07/AutoML_2020_paper_36.pdf
5. Dynamic Pruning for CNNs - Emergent Mind, accessed January 3, 2026, <https://www.emergentmind.com/topics/dynamic-pruning-methodologies-for-cnns>
6. A Closer Look at Sparse Training in Deep Reinforcement Learning - NeurIPS 2025, accessed January 3, 2026, <https://neurips.cc/virtual/2024/109144>
7. dynamic sparse training - arXiv, accessed January 3, 2026, <https://arxiv.org/pdf/2305.02299>
8. Advancing Dynamic Sparse Training by Exploring Optimization Opportunities, accessed January 3, 2026, <https://par.nsf.gov/servlets/purl/10614851>
9. Fisher-Guided Selective Forgetting: Mitigating The Primacy Bias in Deep Reinforcement Learning - arXiv, accessed January 3, 2026, <https://arxiv.org/html/2502.00802v1>
10. On the Variance of the Fisher Information for Deep Learning - NeurIPS, accessed January 3, 2026, <https://proceedings.neurips.cc/paper/2021/file/2d290e496d16c9dcaa9b4ded5cac10cc-Paper.pdf>
11. Behavior Synthesis via Contact-Aware Fisher Information Maximization - arXiv, accessed January 3, 2026, <https://arxiv.org/html/2505.12214v1>
12. Gradient-Weight Alignment as a Train-Time Proxy for Generalization in Classification Tasks, accessed January 3, 2026, <https://arxiv.org/html/2510.25480v1>
13. Unleashing the Power of Gradient Signal-to-Noise Ratio for Zero-Shot NAS - ResearchGate, accessed January 3, 2026, https://www.researchgate.net/publication/377419990_Unleashing_the_Power_of_Gradient_Signal-to-Noise_Ratio_for_Zero-Shot_NAS
14. Unleashing the Power of Gradient Signal-to-Noise Ratio for Zero-Shot NAS - CVF Open Access, accessed January 3, 2026, https://openaccess.thecvf.com/content/ICCV2023/papers/Sun_Unleashing_the_Po

[wer_of_Gradient_Signal-to-Noise_Ratio_for_Zero-Shot_NAS_ICCV_2023_paper.pdf](#)

15. Frozen Language Models Are Gradient Coherence Rectifiers in Vision Transformers, accessed January 3, 2026,
<https://ojs.aaai.org/index.php/AAAI/article/view/32176/34331>
16. L9805E - Super smart power motor driver with 8-Bit MCU, RAM, EEPROM, ADC, WDG, Timer - STMicroelectronics, accessed January 3, 2026,
<https://www.st.com/resource/en/datasheet/l9805e.pdf>
17. Optimization Coaching for JavaScript - DROPS, accessed January 3, 2026,
<https://drops.dagstuhl.de/storage/00lipics/lipics-vol037-ecoop2015/LIPIcs.ECOOP.2015.271/LIPIcs.ECOOP.2015.271.pdf>
18. Neural Graph Embedding for Neural Architecture Search | Request PDF - ResearchGate, accessed January 3, 2026,
https://www.researchgate.net/publication/342540006_Neural_Graph_EMBEDDING_for_Neural_Architecture_Search
19. Neural Graph Embedding for Neural Architecture Search, accessed January 3, 2026, <https://ojs.aaai.org/index.php/AAAI/article/view/5903/5759>
20. Graph Neural Network-Based Surrogate Model for Evolutionary Neural Architecture Search, accessed January 3, 2026,
<https://ieeexplore.ieee.org/document/11214150/>
21. RL-Pruner: Structured Pruning Using Reinforcement Learning for CNN Compression and Acceleration - arXiv, accessed January 3, 2026,
<https://arxiv.org/html/2411.06463v1>
22. Graph Embedding for Neural Architecture Search with Input-Output Information - AutoML Conference 2022, accessed January 3, 2026,
https://2022.automl.cc/wp-content/uploads/2022/07/graph_embedding_for_neural_arc.pdf
23. Efficient Second-Order Methods for Non-Convex Optimization and Machine Learning - UC Berkeley, accessed January 3, 2026,
<https://escholarship.org/content/qt0431q1ws/qt0431q1ws.pdf>
24. A Review of the Gumbel-max Trick and its Extensions for Discrete Stochasticity in Machine Learning - SciSpace, accessed January 3, 2026,
<https://scispace.com/pdf/a-review-of-the-gumbel-max-trick-and-its-extensions-for-13o8upwy.pdf>
25. The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables, accessed January 3, 2026,
<https://vitalab.github.io/article/2018/11/29/concrete.html>
26. Gumbel-Softmax: Differentiable Discrete Sampling - Emergent Mind, accessed January 3, 2026,
<https://www.emergentmind.com/topics/gumbel-softmax-relaxation>
27. Improving Discrete Optimisation Via Decoupled Straight-Through Gumbel-Softmax - arXiv, accessed January 3, 2026,
<https://arxiv.org/html/2410.13331v1>
28. What is Gumbel-Softmax?. A differentiable approximation to... | by Wanshun Wong | TDS Archive | Medium, accessed January 3, 2026,

<https://medium.com/data-science/what-is-gumbel-softmax-7f6d9cdcb90e>

29. Differentiable Subset Pruning of Transformer Heads - ACL Anthology, accessed January 3, 2026, <https://aclanthology.org/2021.tacl-1.86.pdf>
30. Differentiable Subset Pruning of Transformer Heads | Transactions ..., accessed January 3, 2026,
https://direct.mit.edu/tacl/article/doi/10.1162/tacl_a_00436/108868/Differentiable-Subset-Pruning-of-Transformer-Heads
31. Spatial Re-Parameterization for N:M Sparsity - IEEE Computer Society, accessed January 3, 2026,
<https://www.computer.org/csdl/journal/tp/2025/09/11023133/27gSMe44FJm>
32. S-STE: Continuous Pruning Function for Efficient 2:4 Sparse Pre-training - NIPS papers, accessed January 3, 2026,
https://proceedings.neurips.cc/paper_files/paper/2024/file/3b576711b12ab036b45130fc8eb78504-Paper-Conference.pdf
33. MaskPro: Linear-Space Probabilistic Learning for Strict (N:M)-Sparsity on LLMs, accessed January 3, 2026, <https://openreview.net/forum?id=0R06BghLJX>
34. Accelerating GNNs on GPU Sparse Tensor Cores through N:M Sparsity-Oriented Graph Reordering - OSTI.GOV, accessed January 3, 2026,
<https://www.osti.gov/servlets/purl/2524569>
35. Multi-Granular Node Pruning for Circuit Discovery - arXiv, accessed January 3, 2026, <https://arxiv.org/html/2512.10903v1>
36. Efficient Denoising Diffusion via Probabilistic Masking - Yuan Gao's Homepage, accessed January 3, 2026, https://yuan-gao.net/pdf/ICML2024_efficient.pdf
37. Learning Representations in Model-Free Hierarchical Reinforcement Learning, accessed January 3, 2026,
<https://ojs.aaai.org/index.php/AAAI/article/view/5141/5014>
38. Hierarchical Reinforcement Learning on Multi-Channel Hypergraph Neural Network for Course Recommendation - IJCAI, accessed January 3, 2026, <https://www.ijcai.org/proceedings/2024/0232.pdf>
39. A Greedy Hierarchical Approach to Whole-Network Filter- Pruning in CNNs - OpenReview, accessed January 3, 2026,
<https://openreview.net/pdf?id=WzHuebRSgQ>
40. Multi-Agent Actor-Critic with Harmonic Annealing Pruning for Dynamic Spectrum Access Systems - arXiv, accessed January 3, 2026,
<https://arxiv.org/html/2503.15172v1>
41. CTDE: Centralized Training & Decentralized Execution - Emergent Mind, accessed January 3, 2026,
<https://www.emergentmind.com/topics/centralized-training-with-decentralized-execution-ctde>
42. PTDE: Personalized Training with Distilled Execution for Multi-Agent Reinforcement Learning - ResearchGate, accessed January 3, 2026,
https://www.researchgate.net/publication/382788490_PTDE_Personalized_Training_with_Distilled_Execution_for_Multi-Agent_Reinforcement_Learning
43. (PDF) A multi-agent reinforcement learning based approach for ..., accessed January 3, 2026,

https://www.researchgate.net/publication/387498437_A_multi-agent_reinforcement_learning_based_approach_for_automatic_filter_pruning

44. Decentralized Neural Network Policies - Emergent Mind, accessed January 3, 2026,
<https://www.emergentmind.com/topics/decentralized-neural-network-policies>
45. MAT-Agent: Adaptive Multi-Agent Training Optimization - arXiv, accessed January 3, 2026, <https://arxiv.org/html/2510.17845v1>
46. eGAI: Multi-Skilled Policy for Energy-efficient Gait Transitions - UCL Discovery - University College London, accessed January 3, 2026,
<https://discovery.ucl.ac.uk/id/eprint/10217600/1/J27.pdf>
47. Yejie Wang - CatalyzeX, accessed January 3, 2026,
<https://www.catalyzex.com/author/Yejie%20Wang>
48. Explore and Establish Synergistic Effects Between Weight Pruning and Coreset Selection in Neural Network Training - arXiv, accessed January 3, 2026,
<https://arxiv.org/html/2511.09901v1>
49. Continuous action deep reinforcement learning for propofol dosing during general anesthesia | Request PDF - ResearchGate, accessed January 3, 2026,
https://www.researchgate.net/publication/356719364_Continuous_action_deep_reinforcement_learning_for_propofol_dosing_during_general_anesthesia
50. Memory, learning, and cognition (Chapter 15) - Anesthetic Pharmacology, accessed January 3, 2026,
<https://www.cambridge.org/core/books/anesthetic-pharmacology/memory-learning-and-cognition/9C1E9B2182FDB7C84B5A1686245562A1>
51. Multi-Agent Actor-Critic with Harmonic Annealing Pruning for Dynamic Spectrum Access Systems - Eusipco 2025, accessed January 3, 2026,
<https://eusipco2025.org/wp-content/uploads/pdfs/0001937.pdf>
52. 10 Benefits of Using Quantum Computing for Optimization Problems in Aerospace and Defense - BQP, accessed January 3, 2026,
<https://www.bqpsim.com/articles/quantum-optimization-in-aerospace-defense>
53. Neural Network Structure Optimization by Simulated Annealing - PMC, accessed January 3, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8947290/>
54. Prune Neural Network with Simulated Annealing and Genetic Algorithm (Part 1: Overview), accessed January 3, 2026,
https://18520339.medium.com/prune-neural-network-with-simulated-annealing-and-genetic-algorithm-part-1-overview-12b77f6fda0d?source=rss-----artificial_intelligence-5
55. Time-Aware Feature Selection: Adaptive Temporal Masking for Stable Sparse Autoencoder Training - arXiv, accessed January 3, 2026,
<https://arxiv.org/html/2510.08855v1>
56. Time-Aware Feature Selection: Adaptive Temporal Masking for Stable Sparse Autoencoder Training - ResearchGate, accessed January 3, 2026,
https://www.researchgate.net/publication/396457192_Time-Aware_Feature_Selection_Adaptive_Temporal_Masking_for_Stable_Sparse_Autoencoder_Training